

Pharma Validation Library — Expansion Report

AgenThink Mesh | Phase II → Phase III Retrospective Governance Library

10 → 50 Cases | Board-Quality Institutional Diligence Document

Classification: Internal Research & Validation | Not for Distribution

Date: June 2026

Prepared by: AgenThink Mesh Validation Director

Review standard: FDA / EMA / Big Four / Life Sciences PE

EXECUTIVE VERDICT (Opinion First)

At 10 cases: CAUTION — do not commercialise.

At 20 cases: CAUTION — pilot-eligible with documented methodology.

At 50 cases: GO (conditional) — sufficient for enterprise pilots with disclosed limitations.

At 100 cases: GO — peer-review eligible; enterprise deployment defensible.

The Pharma Council has demonstrated genuine signal detection capability. The $\frac{9}{10}$ primary driver identification rate is not a marketing claim — it is a documented finding with one honest failure (Verubecestat) and a traceable root cause. The expansion to 50 cases is necessary before any commercialisation discussion. The 40 proposed cases in this report are selected to stress-test the council against its weakest areas: clean-signal successes (where false-positive blocking is the risk), Alzheimer's amyloid failures (where the council already failed once), and novel mechanism approvals (where the evidence threshold question is most acute).

TABLE OF CONTENTS

1. Deliverable 1 — 40 New Case Proposals
 2. Deliverable 2 — Balanced Validation Set (50-Case Architecture)
 3. Deliverable 3 — Case Database (All 40 Cases, Full Fields)
 4. Deliverable 4 — Validation Metrics Framework
 5. Deliverable 5 — Constitution Learning Framework & Change Log
 6. Deliverable 6 — Statistical Credibility Assessment
 7. Deliverable 7 — Independent Review Readiness
 8. Deliverable 8 — Executive Commercialisation Recommendation
-

DELIVERABLE 1 — 40 New Case Proposals

Selection Methodology

Cases were selected to satisfy four anti-bias criteria simultaneously:

Anti-survivorship bias: The 40 cases include 33 eventual successes, 2 failures, 2 partial successes, 2 regulatory controversies, and 1 ongoing programme. The apparent success-heavy distribution reflects the deliberate inclusion of clean-GO cases — drugs where the council's risk is false-positive blocking, not failure detection. This is the under-tested half of the validation problem.

Anti-selection bias: Cases were drawn from a pre-specified list of therapeutic areas and decision years, not selected after reviewing outcomes. The 10 original cases were selected for governance complexity; the 40 expansion cases are selected for therapeutic area coverage and outcome diversity.

Anti-hindsight bias: All cases have published Phase II data in the peer-reviewed literature or FDA/EMA submission documents. Evidence boundaries can be independently verified.

Anti-overfitting: The expansion deliberately includes 8 immunology cases where the council's constitutional rules were not calibrated — testing whether the rules generalise beyond the original 10 cases.

Case Count by Therapeutic Area

Therapeutic Area	New Cases	Original Cases	Total
Oncology	10	2 (pembrolizumab, entrectinib)	12
Cardiovascular	8	2 (torcetrapib, ticagrelor)	10
Neurology / Alzheimer's	6	3 (aducanumab, verubecestat, solanezumab)	9
Immunology	8	0	8
Metabolic / Diabetes / Obesity	6	1 (evolocumab)	7
Rare Disease / Genetic Medicine	5	2 (eteplirsen, sofosbuvir)	7
Total	43	10	53

Note: Three cases (sacituzumab govitecan, donanemab, nusinersen) were classified under multiple areas by the research engine; they are assigned to their primary area above. Final count: 50 unique cases.

Outcome Distribution (Full 50-Case Library)

Outcome Category	Count	%
Successful	33	66%
Failed	7	14%
Partially Successful	4	8%
Regulatory Controversy	4	8%
Withdrawn / Terminated	2	4%
Total	50	100%

This distribution is intentional. A validation library weighted toward failures would overfit the council's detection rules. A library weighted toward successes tests the equally important false-positive problem. The 66% success rate approximates the historical Phase III success rate for drugs that reach Phase III (~60–65% across therapeutic areas), making this a representative sample.

DELIVERABLE 2 — Balanced Validation Set (50-Case Architecture)

The 50-Case Validation Architecture

The complete library is structured as follows. Cases 1–10 are completed. Cases 11–50 are proposed in this document.

Tier A — Completed Cases (1–10)

#	Drug	Area	Verdict	Outcome	Correct
1	Torcetrapib	Cardiovascular	WAIT	Failed	YES
2	Semagacestat	Neurology	WAIT	Failed	YES
3	Muraglitazar	Cardiovascular	NO-GO	Withdrawn	YES
4	Sofosbuvir	Rare Disease	GO	Successful	YES
5	Aducanumab	Alzheimer's	WAIT	Reg. Controversy	PARTIAL
6	Verubecestat	Alzheimer's	GO	Failed	NO
7	Evolocumab	Cardiovascular	GO	Successful	YES
8	Eteplirsen	Rare Disease	GO	Reg. Controversy	PARTIAL
9	Ticagrelor	Cardiovascular	WAIT	Successful	YES
10	Entrectinib	Oncology	GO	Successful	YES

Tier B — Proposed Oncology Cases (11–20)

#	Drug	Company	Phase II Year	Outcome	Governance Test
11	Pembrolizumab	Merck	2014	Successful	Accelerated approval on ORR — surrogate-to-survival translation
12	Nivolumab	BMS	2013	Successful	First PD-1 inhibitor — novel mechanism, small Phase II
13	Olaparib	AstraZeneca	2012	Successful	BRCA biomarker selection — precision medicine evidence threshold
14	Osimertinib	AstraZeneca	2015	Successful	Third-generation EGFR — resistance mechanism evidence
15	Venetoclax	AbbVie/Roche	2014	Successful	BCL-2 inhibitor — tumour lysis syndrome safety management
16	Ribociclib	Novartis	2015	Successful	CDK4/6 inhibitor — QTc prolongation signal management
17	Durvalumab	AstraZeneca	2015	Successful	Stage III NSCLC — unresectable disease, new treatment paradigm
18	Lorlatinib	Pfizer	2017	Successful	Third-generation ALK — CNS penetration evidence
19	Sacituzumab govitecan	Immunomedics	2017	Successful	ADC in TNBC — accelerated approval on single-arm data
20	Adagrasib	Mirati	2020	Successful	KRAS G12C — first-in-class, accelerated approval

Tier C — Proposed Cardiovascular Cases (21–28)

#	Drug	Company	Phase II Year	Outcome	Governance Test
21	Alirocumab	Sanofi/Regeneron	2013	Successful	PCSK9 inhibitor — LDL-C surrogate vs. CV outcomes
22	Inclisiran	Amylin/Novartis	2017	Successful	siRNA mechanism — twice-yearly dosing, adherence assumptions
23	Sacubitril/valsartan	Novartis	2012	Successful	PARADIGM-HF — superiority over ACE inhibitor, stopping early
24	Rivaroxaban	Bayer/J&J	2007	Successful	ROCKET AF — non-inferiority design, warfarin comparator
25	Cangrelor	Medicines Co.	2008	Successful	Two Phase III failures before eventual approval — persistence
26	Vorapaxar	Merck	2008	Partially Successful	Approved with black box warning — benefit/risk in subgroups
27	Mavacamten	MyoKardia	2019	Successful	HCM — novel mechanism, LVOT obstruction endpoint
28	Selatogrel	Idorsia	2020	Ongoing	Self-administered SC antiplatelet — Phase III ongoing

Tier D — Proposed Neurology / Alzheimer's Cases (29–34)

#	Drug	Company	Phase II Year	Outcome	Governance Test
29	Lecanemab	Eisai/Biogen	2020	Successful	Amyloid antibody — modest clinical benefit, ARIA safety
30	Donanemab	Eli Lilly	2021	Successful	Tau-PET staging — novel patient selection
31	Gantenerumab	Roche	2015	Failed	Phase III failure 2022 — amyloid reduction without clinical benefit
32	Crenezumab	Genentech	2015	Failed	Phase III failure 2019 — dose selection question
33	Brexanolone	Sage Therapeutics	2016	Reg. Controversy	PPD — novel GABA mechanism, IV administration barrier
34	Zuranolone	Sage/Biogen	2020	Reg. Controversy	MDD — mixed Phase III, FDA declined MDD indication

Tier E — Proposed Immunology Cases (35–42)

#	Drug	Company	Phase II Year	Outcome	Governance Test
35	Dupilumab	Sanofi/Regeneron	2014	Successful	IL- γ_{13} dual blockade — atopic dermatitis, new mechanism
36	Secukinumab	Novartis	2012	Successful	IL-17A — psoriasis, neutropenia signal management
37	Ixekizumab	Eli Lilly	2013	Successful	IL-17A — crowded class, differentiation evidence
38	Guselkumab	Janssen	2015	Successful	IL-23 p19 — selective vs. broad IL-23 inhibition
39	Risankizumab	AbbVie	2016	Successful	IL-23 p19 — Crohn's extension from psoriasis data
40	Ozanimod	BMS	2016	Successful	S1P receptor — MS and UC, cardiac monitoring
41	Upadacitinib	AbbVie	2016	Successful	JAK1 selective — RA, VTE and malignancy class risk
42	Bimekizumab	UCB	2018	Successful	IL-17A/F dual — differentiation from IL-17A only

Tier F — Proposed Metabolic / Diabetes / Obesity Cases (43–48)

#	Drug	Company	Phase II Year	Outcome	Governance Test
43	Semaglutide	Novo Nordisk	2017	Successful	GLP-1 — obesity indication, CV outcomes, dose escalation
44	Tirzepatide	Eli Lilly	2019	Successful	GIP/GLP-1 dual — unprecedented weight loss, class differentiation
45	Dapagliflozin	AstraZeneca	2009	Successful	SGLT2 — DKA signal, CV/renal benefits not anticipated
46	Empagliflozin	BI/Lilly	2010	Successful	SGLT2 — EMPA-REG CV mortality surprise
47	Lixisenatide	Sanofi	2011	Partially Successful	GLP-1 — ELIXA non-inferiority only, no CV superiority
48	Canagliflozin	J&J	2011	Partially Successful	SGLT2 — amputation signal in CANVAS, regulatory action

Tier G — Proposed Rare Disease / Genetic Medicine Cases (49–50 + 3 from original)

#	Drug	Company	Phase II Year	Outcome	Governance Test
49	Patisiran	Alnylam	2015	Successful	First RNAi — hATTR, novel mechanism, LNP delivery
50	Onasemnogene abeparvovec	AveXis/Novartis	2018	Successful	Gene therapy SMA1 — \$2.1M price, single-dose curative intent
(ref)	Voretigene neparvovec	Spark	2015	Successful	First in vivo gene therapy — RPE65, n=31 pivotal trial

DELIVERABLE 3 — Case Database (All 40 Proposed Cases, Full Fields)

The following table presents all 40 proposed expansion cases with the full field set required for validation library inclusion. Cases are sorted by therapeutic area and ranking score.

Oncology (10 Cases)

Field	Pembrolizumab	Nivolumab	Olaparib	Osimertinib	Venetoclax
Brand	Keytruda	Opdivo	Lynparza	Tagrisso	Venclexta
Company	Merck	BMS	AstraZeneca	AstraZeneca	AbbVie/Roche
Indication	Melanoma / NSCLC	Melanoma	BRCA+ ovarian	EGFR T790M NSCLC	CLL / AML
Phase II Year	2014	2013	2012	2015	2014
Outcome	Successful	Successful	Successful	Successful	Successful
Regulatory	FDA 2014 (acc.)	FDA 2014 (acc.)	FDA 2014 (acc.)	FDA 2015 (acc.)	FDA 2016
Financial Impact	>\$25B peak sales	>\$9B peak sales	>\$3B peak sales	>\$5B peak sales	>\$2B peak sales
Value at Risk	\$1–2B	\$1–2B	\$800M	\$1B	\$1B
Evidence Availability	High	High	High	High	High
Ranking Score	9	9	9	9	9
Inclusion	Yes	Yes	Yes	Yes	Yes

Field	Ribociclib	Durvalumab	Lorlatinib	Sacituzumab govitecan	Adagrasib
Brand	Kisqali	Imfinzi	Lorbrena	Trodelvy	Krazati
Company	Novartis	AstraZeneca	Pfizer	Immunomedics	Mirati
Indication	HR+/HER2- BC	Stage III NSCLC	ALK+ NSCLC	TNBC	KRAS G12C NSCLC
Phase II Year	2015	2015	2017	2017	2020
Outcome	Successful	Successful	Successful	Successful	Successful
Regulatory	FDA 2017	FDA 2018	FDA 2018	FDA 2020 (acc.)	FDA 2022 (acc.)
Financial Impact	>\$2B peak	>\$3B peak	>\$500M peak	>\$1B peak	>\$500M peak
Value at Risk	\$800M	\$1B	\$600M	\$500M	\$400M
Evidence Availability	High	High	High	High	High
Ranking Score	9	9	9	9	9
Inclusion	Yes	Yes	Yes	Yes	Yes

Key governance tests — Oncology: Pembrolizumab and nivolumab test the council's ability to issue a GO on accelerated approval based on objective response rate (ORR) without overall survival data — the mirror image of the Aducanumab problem. Olaparib tests precision medicine evidence thresholds: BRCA biomarker selection with a small Phase II. Venetoclax tests safety management: tumour lysis syndrome was a known and manageable risk. Ribociclib tests signal proportionality: QTc prolongation was present in Phase II but the council must weigh it against the efficacy signal in a high-unmet-need population.

Cardiovascular (8 Cases)

Field	Alirocumab	Inclisiran	Sacubitril/valsartan	Rivaroxaban	Cangrelor
Brand	Praluent	Leqvio	Entresto	Xarelto	Kengreal
Company	Sanofi/Regeneron	Alynham/Novartis	Novartis	Bayer/J&J	Medicines Co.
Indication	Hypercholesterolemia	Hypercholesterolemia	Heart failure	AF/VTE	ACS/PCI
Phase II Year	2013	2017	2012	2007	2008
Outcome	Successful	Successful	Successful	Successful	Successful
Regulatory	FDA 2015	FDA 2021	FDA 2015	FDA 2011	FDA 2015
Financial Impact	~\$600M peak	>\$1B projected	>\$4B peak	>\$6B peak	~\$100M peak
Value at Risk	\$2B	\$1B	\$2B	\$3B	\$500M
Evidence Availability	High	High	High	High	High
Ranking Score	8	9	9	9	8
Inclusion	Yes	Yes	Yes	Yes	Yes

Field	Vorapaxar	Mavacamten	Selatogrel
Brand	Zontivity	Camzyos	N/A
Company	Merck	MyoKardia	Idorsia
Indication	Secondary ACS prevention	Hypertrophic cardiomyopathy	Acute MI
Phase II Year	2008	2019	2020
Outcome	Partially Successful	Successful	Ongoing
Regulatory	FDA 2014 (black box)	FDA 2022	Phase III
Financial Impact	<\$100M peak (restricted)	>\$500M projected	Unknown
Value at Risk	\$2B	\$1B	\$500M
Evidence Availability	High	High	High
Ranking Score	8	9	8
Inclusion	Yes	Yes	Yes

Key governance tests — Cardiovascular: Cangrelor is the most instructive: two Phase III failures (CHAMPION-PCI, CHAMPION-PLATFORM) before eventual approval in CHAMPION-PHOENIX. The council would have encountered the first two failures — testing whether it correctly identifies the endpoint selection problem rather than the drug itself. Vorapaxar tests benefit/risk in subgroups: the drug works in post-MI patients without prior stroke but causes fatal bleeding in stroke patients. Sacubitril/valsartan tests early stopping: PARADIGM-HF was stopped early for benefit — the council must evaluate whether early stopping inflates the effect estimate.

Neurology / Alzheimer’s (6 Cases)

Field	Lecanemab	Donanemab	Gantenerumab	Crenezumab	Brexanolone	Zuranolone
Brand	Leqembi	N/A	N/A	N/A	Zulresso	Zurzuvae
Company	Eisai/Biogen	Eli Lilly	Roche	Genentech	Sage	Sage/Biogen
Indication	Early AD	Early AD	Early AD	AD	PPD	MDD
Phase II Year	2020	2021	2015	2015	2016	2020
Outcome	Successful	Successful	Failed	Failed	Reg. Controversy	Reg. Controversy
Regulatory	FDA 2023 (acc.)	FDA 2024	Terminated 2022	Terminated 2019	FDA 2019	FDA 2023 (PPD only)
Financial Impact	>\$1B projected	>\$2B projected	~\$1B write-off	~\$500M write-off	~\$100M peak	~\$250M peak
Value at Risk	\$2B	\$3B	\$2B	\$1B	\$500M	\$1B
Evidence Availability	High	High	High	High	High	High
Ranking Score	9	9	9	8	8	8
Inclusion	Yes	Yes	Yes	Yes	Yes	Yes

Key governance tests — Neurology / Alzheimer’s: Gantenerumab and crenezumab are the most important cases in this tier. Both showed amyloid plaque reduction in Phase II. Both failed Phase III. These cases test whether the council correctly identifies the surrogate endpoint problem in Alzheimer’s — the same constitutional gap that caused the Verubecestat failure. If the council issues GO on gantenerumab and crenezumab after the proposed PC-011 amendment, the amendment has not solved the problem. If it issues WAIT, the amendment is validated. Lecanemab and donanemab test the opposite: modest but real clinical benefit — the council must not block drugs that work.

Immunology (8 Cases)

Field	Dupilumab	Secukinumab	Ixekizumab	Guselkumab	Risankizumab	Ozanimod	Upadacitinib	Bimek
Brand	Dupixent	Cosentyx	Taltz	Tremfya	Skyrizi	Zeposia	Rinvoq	Bimzi
Company	Sanofi/Regen.	Novartis	Eli Lilly	Janssen	AbbVie	BMS	AbbVie	UCB
Indication	Atopic dermatitis	Psoriasis/PsA/AS	Psoriasis	Psoriasis	Psoriasis/CD	MS / UC	RA	Psori
Phase II Year	2014	2012	2013	2015	2016	2016	2016	2018
Outcome	Successful	Successful	Successful	Successful	Successful	Successful	Successful	Succo
Regulatory	FDA 2017	FDA 2015	FDA 2016	FDA 2017	FDA 2019	FDA 2020	FDA 2019	FDA :
Financial Impact	\$14.9B (2024)	>\$5B peak	>\$2B peak	>\$2B peak	>\$4B peak	>\$1B peak	>\$4B peak	>\$1B proje
Evidence Availability	High	High	High	High	High	High	High	High
Ranking Score	9	9	9	9	9	9	9	9
Inclusion	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes

Key governance tests — Immunology: The immunology tier is the most important test of false-positive blocking. All 8 drugs succeeded. If the council issues WAIT or NO-GO on any of them, it is generating false positives — the most commercially damaging error for a governance tool. Upadacitinib tests class-level safety: JAK inhibitors carry VTE and malignancy risks that were known from tofacitinib. The council must weigh class risk against drug-specific Phase II data. Ozanimod tests cardiac monitoring requirements: S1P receptor modulators require first-dose cardiac observation — a known class effect that should not block advancement.

Metabolic / Diabetes / Obesity (6 Cases)

Field	Semaglutide	Tirzepatide	Dapagliflozin	Empagliflozin	Lixisenatide	Canagliflozin
Brand	Ozempic/Wegovy	Mounjaro/Zepbound	Farxiga	Jardiance	Adlyxin	Invokana
Company	Novo Nordisk	Eli Lilly	AstraZeneca	BI/Lilly	Sanofi	J&J
Indication	Obesity / T2D	T2D / Obesity	T2D	T2D	T2D + CV	T2D
Phase II Year	2017	2019	2009	2010	2011	2011
Outcome	Successful	Successful	Successful	Successful	Partially Successful	Partially Successful
Regulatory	FDA ²⁰¹⁷ / ₂₀₂₁	FDA ²⁰²² / ₂₀₂₃	FDA 2014	FDA 2014	FDA 2016	FDA 2013 (later restricted)
Financial Impact	>\$20B projected	>\$15B projected	>\$3B peak	>\$5B peak	<\$200M peak	>\$1B peak (then declined)
Evidence Availability	High	High	High	High	High	High
Ranking Score	9	10	9	10	7	8
Inclusion	Yes	Yes	Yes	Yes	Yes	Yes

Key governance tests — Metabolic / Diabetes / Obesity: Dapagliflozin and empagliflozin are the most instructive cases. Both were approved on glucose-lowering efficacy. The CV and renal benefits were not anticipated at the Phase II decision point. The council must evaluate these drugs on the evidence available at the time — not with knowledge of EMPA-REG OUTCOME (2015) or DECLARE-TIMI 58 (2018). Tirzepatide tests unprecedented efficacy: ~22% body weight reduction in Phase II was beyond any prior GLP-1 data. The council must evaluate whether extraordinary efficacy claims require extraordinary evidence. Canagliflozin tests post-approval signal management: the amputation signal in CANVAS was not visible in Phase II but was a class-level concern.

Rare Disease / Genetic Medicine (5 Cases)

Field	Patisiran	Onasemnogene abeparvovec	Voretigene neparvovec	Nusinersen	Givosiran
Brand	Onpattro	Zolgensma	Luxturna	Spinraza	Givlaari
Company	Alnylam	AveXis/Novartis	Spark	Biogen/Ionis	Alnylam
Indication	hATTR amyloidosis	SMA type 1	RPE65 retinal dystrophy	SMA	Acute hepatic porphyria
Phase II Year	2015	2018	2015	2015	2018
Outcome	Successful	Successful	Successful	Successful	Successful
Regulatory	FDA 2018	FDA 2019	FDA 2017	FDA 2016	FDA 2019
Financial Impact	>\$500M peak	>\$1B peak	~\$400M peak	>\$2B peak	~\$200M peak
Value at Risk	\$1B	\$2B	\$500M	\$1B	\$500M
Evidence Availability	High	High	Medium	High	High
Ranking Score	9	9	8	9	8
Inclusion	Yes	Yes	Yes	Yes	Yes

Key governance tests — Rare Disease / Genetic Medicine: Onasemnogene abeparvovec tests the extreme end of the evidence threshold: a single-arm Phase I/II study in 15 patients was the basis for FDA approval. The \$2.1M list price makes this the highest-stakes pricing governance question in the library. Patisiran tests first-in-class mechanism acceptance: RNAi therapeutics had never been approved; the council must evaluate a genuinely novel delivery mechanism. Voretigene neparvovec tests the smallest pivotal dataset: n=31 in the pivotal trial. The council must distinguish between “insufficient evidence” and “maximum feasible evidence in an ultra-rare disease.”

DELIVERABLE 4 — Validation Metrics Framework

Current Metrics (Carried Forward)

The 10-case library uses five metrics. These are preserved and extended:

Metric	Definition	10-Case Value
Verdict Alignment (strict)	% cases where council verdict matches actual outcome	70% ($\frac{7}{10}$)
Verdict Alignment (broad)	% cases where council identified primary decision driver	90% ($\frac{9}{10}$)
Overall Validation Score (OVS)	Composite score per case (0–10)	8.73 avg
False Positive Rate	% successful drugs blocked by council	0% (0/5)
False Negative Rate	% failed drugs not caught by council	20% ($\frac{1}{5}$)

Extended Institutional Metrics

The following nine metrics are added for the 50-case library. Exact formulas are provided.

1. Sensitivity (True Positive Rate)

Definition: Of all drugs that actually failed Phase III or were withdrawn, what fraction did the council correctly flag as WAIT or NO-GO?

Formula:

$$\text{Sensitivity} = \text{TP} / (\text{TP} + \text{FN})$$

Where:

- TP = council issued WAIT/NO-GO AND drug failed/was withdrawn
- FN = council issued GO AND drug failed/was withdrawn

10-case value: $\frac{4}{5} = 0.80$ (Verubecestat is the FN)

Target at 50 cases: ≥ 0.75 (allowing for 3–4 genuine failures in a 50-case library)

2. Specificity (True Negative Rate)

Definition: Of all drugs that actually succeeded, what fraction did the council correctly allow through (GO or WAIT with eventual GO)?

Formula:

$$\text{Specificity} = \text{TN} / (\text{TN} + \text{FP})$$

Where:

- TN = council issued GO AND drug succeeded
- FP = council issued WAIT/NO-GO AND drug succeeded

10-case value: $\frac{5}{5} = 1.00$ (no false positives in original library)

Expected at 50 cases: Will decrease as clean-success cases are added. Target: ≥ 0.70

Critical note: The 1.00 specificity in the original library is an artefact of case selection — all 5 success cases had complex governance questions. The 40 expansion cases include 33 clean successes that will test this metric properly.

3. Precision (Positive Predictive Value)

Definition: Of all cases where the council issued WAIT or NO-GO, what fraction actually failed or were withdrawn?

Formula:

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

10-case value: $\frac{4}{4} = 1.00$ (all WAIT/NO-GO verdicts were retrospectively correct)

Expected at 50 cases: Will decrease. Target: ≥ 0.60

4. Recall

Definition: Identical to Sensitivity. Included separately for alignment with standard ML evaluation frameworks used by Big Four diligence teams.

Formula:

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

10-case value: 0.80

5. Balanced Accuracy

Definition: Average of Sensitivity and Specificity. Corrects for class imbalance (more successes than failures in the library).

Formula:

$$\text{Balanced Accuracy} = (\text{Sensitivity} + \text{Specificity}) / 2$$

10-case value: $(0.80 + 1.00) / 2 = 0.90$

Target at 50 cases: ≥ 0.75

Note: Balanced Accuracy is the primary metric for regulatory review readiness. A value ≥ 0.75 at 50 cases would be a meaningful finding.

6. Decision Driver Identification Accuracy (DDIA)

Definition: Of all completed cases, what fraction had the council correctly identify the primary decision driver (the factor that most influenced the actual outcome)?

Formula:

$$\text{DDIA} = \text{Cases where primary driver correctly identified} / \text{Total cases}$$

Assessment method: Independent pharmaceutical expert review of council deliberation records vs. published post-outcome analyses.

10-case value: $\frac{9}{10} = 0.90$ (per internal assessment — requires independent validation)

Target at 50 cases: ≥ 0.80 (with independent expert validation)

7. Constitutional Coverage Score (CCS)

Definition: For each case, what fraction of the material risk factors were covered by at least one constitutional rule?

Formula:

$$\text{CCS} = \sum (\text{Rules triggered per case} / \text{Material risk factors per case}) / \text{Total cases}$$

10-case value: Not formally calculated — requires retroactive audit

Target at 50 cases: ≥ 0.75 (meaning 75% of material risks are covered by existing constitutional rules)

8. Amendment Frequency

Definition: Number of constitutional amendments proposed per 10 completed cases.

Formula:

$$\text{Amendment Frequency} = \text{Amendments proposed} / (\text{Cases completed} / 10)$$

10-case value: 1 amendment (PC-011) = **1.0 per 10 cases**

Target: ≤ 2 per 10 cases (high amendment frequency indicates constitutional instability)

Concern threshold: >3 per 10 cases (suggests the constitution is being retrofitted to failures rather than genuinely learning)

9. Therapeutic Area Robustness Score (TARS)

Definition: Balanced Accuracy calculated separately for each therapeutic area. Identifies areas where the council underperforms.

Formula:

$$\text{TARS}(\text{area}) = \text{Balanced Accuracy for cases in that therapeutic area}$$

10-case values (provisional):

Area	Cases	Sensitivity	Specificity	TARS
Cardiovascular	3	1.00	1.00	1.00
Neurology/AD	3	0.50	1.00	0.75
Rare Disease	2	1.00	1.00	1.00
Oncology	1	N/A	1.00	N/A
Metabolic	1	N/A	1.00	N/A

Critical finding: Neurology/Alzheimer's has the lowest TARS (0.75) with only 3 cases. The 6 proposed expansion cases in this area will either confirm or refute this finding.

Metrics Dashboard (Target State at 50 Cases)

Metric	10-Case	Target at 50 Cases	Minimum Acceptable
Sensitivity	0.80	≥0.75	0.65
Specificity	1.00	≥0.70	0.60
Precision	1.00	≥0.60	0.50
Balanced Accuracy	0.90	≥0.75	0.65
DDIA	0.90*	≥0.80	0.70
CCS	Not calculated	≥0.75	0.60
Amendment Frequency	1.0/10	≤2.0/10	≤3.0/10

*Requires independent validation

DELIVERABLE 5 — Constitution Learning Framework & Change Log

The Learning Protocol

After every 10 completed cases, the following review is mandatory:

- Step 1 — Gap Identification:** For every case where the council issued an incorrect verdict or missed a material risk, identify which constitutional rule (if any) should have caught it.
- Step 2 — Blind Spot Mapping:** Across all 10 cases, identify recurring patterns in what the council missed. A pattern appearing in 2+ cases constitutes a systemic gap, not a case-specific failure.
- Step 3 — Over-Restriction Audit:** For every case where the council issued WAIT or NO-GO on a drug that succeeded, identify whether the objection was proportionate to the evidence. Over-restriction is as damaging as under-restriction.
- Step 4 — Amendment Drafting:** For each identified gap, draft a proposed constitutional amendment with: rationale, expected impact, historical cases affected, and risk of over-restriction.
- Step 5 — Amendment Freeze:** No amendment takes effect until it has been tested against at least 5 historical cases where it would have applied. This prevents retrofitting.

Pharma Constitution Change Log

Version 1.0 — Original (Cases 1–10)

Effective: Cases 1–10
Rules: PC-001 through PC-010
Status: Active

Rule	Description	Cases Where Triggered
PC-001	Phase II safety signal investigation requirement	Torcetrapib (BP), Muraglitazar (CV)
PC-002	Surrogate endpoint validation requirement	Aducanumab (amyloid), Evolocumab (LDL-C)
PC-003	Financial pressure disclosure requirement	Torcetrapib (\$800M), Aducanumab
PC-004	DSMB charter adequacy requirement	Torcetrapib, Semagacestat
PC-005	Competitive landscape differentiation requirement	Ticagrelor (vs. clopidogrel)
PC-006	Regulatory pathway clarity requirement	Eteplirsen (accelerated approval)
PC-007	Patient population definition requirement	Aducanumab (early vs. late AD)
PC-008	Dose selection justification requirement	Semagacestat, Verubecestat
PC-009	Mechanism plausibility requirement	Sofosbuvir (NS5B polymerase)
PC-010	Evidence completeness threshold	All cases

Version 1.1 — Proposed Amendment (Post-Case 6: Verubecestat)

Trigger: Verubecestat failure — council issued GO; drug failed Phase III for futility

Root cause: No constitutional rule specifically addressed the Alzheimer's disease evidence threshold problem: the disconnect between amyloid/tau biomarker reduction and clinical benefit

Proposed Amendment:

PC-011 — Alzheimer's Disease and Neurodegeneration Evidence Threshold

For any drug targeting a neurodegenerative disease (Alzheimer's disease, Parkinson's disease, frontotemporal dementia, or related conditions), the council shall require:

- (a) Evidence that the proposed biomarker endpoint (amyloid PET, tau PET, CSF biomarkers, or equivalent) has been validated as a surrogate for clinical benefit in at least one completed Phase III trial in the same indication; OR
- (b) A pre-specified clinical benefit endpoint (CDR-SB, ADAS-Cog, MMSE, or equivalent) demonstrating statistically significant improvement in Phase II; OR

© A documented rationale from the sponsor explaining why (a) and (b) cannot be satisfied, with independent expert review.

Absence of (a), (b), or © shall constitute a WAIT condition, not a blocking condition, pending resolution.

Expected impact: Would have correctly flagged Verubecestat, Gantenerumab, and Crenezumab as WAIT. Would not have blocked Lecanemab or Donanemab (both showed clinical benefit on CDR-SB).

Historical cases affected: Verubecestat (case 6), Aducanumab (case 5), proposed cases 31 (Gantenerumab), 32 (Crenezumab), 29 (Lecanemab), 30 (Donanemab)

Risk of over-restriction: Low. PC-011 requires WAIT, not NO-GO. It does not block advancement; it requires resolution of the evidence gap.

Validation test: Apply PC-011 retroactively to cases 29–32. Expected result: WAIT on Gantenerumab and Crenezumab (correct); GO on Lecanemab and Donanemab (correct).

Status: Proposed. Not yet ratified. Requires 5-case validation test before activation.

Version 1.2 — Anticipated Amendment (Post-Cases 11–20: Oncology)

Anticipated trigger: Accelerated approval cases (pembrolizumab, nivolumab, sacituzumab govitecan, adagrasib) — all approved on ORR without OS data

Anticipated gap: No constitutional rule addresses the accelerated approval pathway specifically. The council may issue WAIT on drugs that are appropriate for accelerated approval, generating false positives.

Anticipated amendment:

PC-012 — Accelerated Approval Pathway Recognition

For drugs targeting serious or life-threatening conditions with unmet medical need, the council shall evaluate the accelerated approval pathway as a legitimate regulatory route. A drug with compelling Phase II ORR data in a population with no approved

therapy shall not receive a WAIT verdict solely on the basis of absent Phase III OS data, provided: (a) the sponsor has committed to a confirmatory Phase III trial; and (b) the FDA/EMA has indicated willingness to consider accelerated approval.

Status: Anticipated. To be drafted after oncology cases are completed.

Version 1.3 — Anticipated Amendment (Post-Cases 35–42: Immunology)

Anticipated trigger: JAK inhibitor class risk cases (upadacitinib) — class-level safety signals from tofacitinib may cause false-positive WAIT verdicts on superior drugs

Anticipated gap: No constitutional rule distinguishes between class-level safety signals and drug-specific safety signals.

Anticipated amendment:

PC-013 — Class-Level vs. Drug-Specific Safety Signal Distinction

When a safety signal is attributable to a drug class rather than the specific molecule under review, the council shall require evidence that the specific drug shares the class-level mechanism of toxicity before applying the signal as a blocking condition.

Status: Anticipated. To be drafted after immunology cases are completed.

DELIVERABLE 6 — Statistical Credibility Assessment

The Honest Assessment

The following analysis applies standard statistical reasoning to the validation library. It does not optimise for positive conclusions.

At 10 Cases — Exploratory Findings Only

What can be stated:

- The council correctly identified the primary decision driver in 9 of 10 cases (90%). This is a descriptive finding, not a statistical claim.
- The false positive rate is 0% across 5 success cases. This is insufficient to estimate a population false positive rate.
- The false negative rate is 20% ($\frac{1}{5}$ failure cases missed). With $n=5$, the 95% confidence interval for the true false negative rate is approximately 1%–72% (Wilson interval). This range is too wide to support any claim.

What cannot be stated:

- Any accuracy rate with confidence intervals
- Any comparison to existing pharmaceutical decision-making processes
- Any claim of statistical significance
- Any generalisation to untested therapeutic areas

Appropriate language: “Proof of concept. The council demonstrated signal detection capability in 10 historical cases. No statistical claims are warranted.”

At 20 Cases — Early Validation

What can be stated:

- Point estimates for sensitivity, specificity, and balanced accuracy with wide confidence intervals
- Comparison to the 50% base rate (coin flip) — if balanced accuracy exceeds 0.70, this is statistically distinguishable from chance at $p<0.05$ with $n=20$ (one-sample binomial test)
- Descriptive statistics by therapeutic area (2–4 cases per area)

What cannot be stated:

- Sensitivity/specificity claims with narrow confidence intervals
- Therapeutic area-specific performance claims
- Peer-review-quality findings

Appropriate language: “Early validation. Point estimates suggest [X]% balanced accuracy (95% CI: [Y]–[Z]%). Sample size is insufficient for definitive conclusions.”

Minimum sample size for statistical significance vs. 50% base rate: Using a one-sided binomial test at $\alpha=0.05$, power=0.80:

- To detect 75% balanced accuracy vs. 50%: $n=22$
- To detect 80% balanced accuracy vs. 50%: $n=16$
- To detect 70% balanced accuracy vs. 50%: $n=35$

At 20 cases, AgenThink can claim statistical significance vs. random chance if balanced accuracy $\geq 75\%$.

At 50 Cases — Meaningful Validation

What can be stated:

- Sensitivity and specificity with 95% confidence intervals of ± 10 –15 percentage points
- Therapeutic area-specific performance (5–10 cases per area)
- Comparison to published base rates for Phase III success prediction
- Amendment frequency and constitutional stability metrics
- Preliminary peer-review submission (methods paper)

What cannot be stated:

- Superiority to specific existing processes (no comparator arm)
- Regulatory submission use
- Prospective accuracy claims

Appropriate language: “Meaningful validation. The council demonstrated [X]% balanced accuracy (95% CI: [Y]–[Z]%) across 50 retrospective cases spanning [N] therapeutic areas. One constitutional amendment was required after [N] cases. These findings support enterprise pilot evaluation.”

Statistical power at 50 cases:

- 95% CI for sensitivity: ± 14 percentage points (assuming ~15 failure cases)
- 95% CI for specificity: ± 10 percentage points (assuming ~35 success cases)
- Minimum detectable difference vs. 50% base rate: 15 percentage points ($\alpha=0.05$, power=0.80)

At 100 Cases — Statistically Credible Evidence

What can be stated:

- Sensitivity and specificity with 95% confidence intervals of ± 8 –10 percentage points
- Therapeutic area-specific performance with statistical significance
- Comparison to published Phase III success prediction models (e.g., Hay et al. 2014)
- Constitutional stability assessment (amendment frequency over time)
- Peer-review submission with sufficient power for inferential statistics
- Enterprise deployment with disclosed limitations

What cannot be stated:

- Regulatory submission use without prospective validation
- Superiority to human expert review without a controlled comparison study

Appropriate language: “Statistically credible evidence. The council demonstrated [X]% balanced accuracy (95% CI: [Y]–[Z]%) across 100 retrospective cases. This finding is statistically significant vs. random chance ($p<0.001$) and comparable to published Phase III success prediction benchmarks.”

Summary Table

Sample Size	Claim Level	Balanced Accuracy CI	Appropriate Use
10	Proof of concept	±30 pp	Internal demonstration only
20	Early validation	±20 pp	Pilot discussions with disclosed limitations
50	Meaningful validation	±12 pp	Enterprise pilot evaluation
100	Statistically credible	±8 pp	Peer review; enterprise deployment
200+	Regulatory-grade	±5 pp	Regulatory submission consideration

pp = percentage points

DELIVERABLE 7 — Independent Review Readiness

Reviewer Profiles and Primary Concerns

FDA Reviewer

Primary concern: Is the evidence boundary verifiable? Can an independent reviewer confirm that no post-outcome information was used in the council deliberation?

Strongest evidence: The torcetrapib case has a verifiable evidence boundary (NEJM 2004 BP signal published; ILLUMINATE terminated December 2006). The boundary is independently checkable.

Weakest evidence: The LLM-based council cannot provide a verifiable audit trail showing that its training data did not include post-outcome information. This is the fundamental unfalsifiable claim.

Credibility risk: High. The FDA will ask: “How do you know the model didn’t know the answer?” The honest answer is: “We cannot prove it didn’t.” This must be disclosed.

Remediation: Build a document-retrieval architecture where the council’s input is a curated corpus of pre-cutoff documents (DOI-cited, date-stamped). The council’s output can then be audited against a known input set.

Reviewer objection: “Your evidence boundary is a prompt instruction, not a verifiable constraint. I can’t audit a prompt.”

EMA Reviewer

Primary concern: Is the methodology reproducible? Can a different team run the same cases and get the same results?

Strongest evidence: The constitutional rules are published. The persona definitions are published. The decision brief format is standardised.

Weakest evidence: The exact prompts, model version, temperature settings, and system configuration are not published. Two runs of the same case with the same inputs may produce different outputs.

Credibility risk: High. EMA requires reproducibility as a prerequisite for any regulatory use.

Remediation: Publish the exact methodology: model version, temperature, system prompt, constitutional rules, persona definitions, decision brief template. Run each case twice and report inter-run consistency.

Reviewer objection: “I ran your torcetrapib case with your published methodology and got a different vote distribution. Which result is correct?”

Top 20 Pharma Company

Primary concern: Does this add value over existing processes? What does it catch that our internal teams miss?

Strongest evidence: The torcetrapib case. Pfizer’s internal teams missed the BP signal. The council caught it. This is a concrete, documented example of value-add.

Weakest evidence: The immunology cases. All 8 drugs succeeded. The council will likely issue GO on all 8. This demonstrates the council doesn’t block good drugs — but it doesn’t demonstrate it adds value over a competent internal team.

Credibility risk: Medium. The pharma company will ask for a prospective pilot, not a retrospective library.

Remediation: Offer a blinded prospective pilot: provide the council with current Phase II data for 5 drugs with unknown outcomes. Report results when Phase III data becomes available (2–3 year timeline).

Reviewer objection: “Your retrospective library is interesting, but I already know the outcomes. Show me what you'd say about a drug I'm deciding on today.”

Big Four Advisory Team

Primary concern: Is the methodology defensible in a diligence report? Can we cite this in an IC memo?

Strongest evidence: The documented failure (Verubecestat) with a traceable root cause and a proposed constitutional amendment. This demonstrates intellectual honesty.

Weakest evidence: The 0% false positive rate across 5 success cases. A Big Four team will immediately note that this is based on a hand-selected sample.

Credibility risk: Medium. Big Four teams are sophisticated consumers of methodology disclosures.

Remediation: Publish the selection criteria for the original 10 cases. Demonstrate that the 40 expansion cases were selected before outcomes were reviewed (pre-specification).

Reviewer objection: “You selected 10 cases that included 5 famous failures and 5 famous successes. Of course you got good results. Show me the cases you didn't include.”

Life Sciences PE Fund

Primary concern: What is the commercial value? How does this translate to a defensible investment thesis?

Strongest evidence: The torcetrapib case. 800*Mwrite – off*+21B market cap loss. If the council had been used, the GO decision would have been challenged. The value of a WAIT verdict in that case is quantifiable.

Weakest evidence: The lack of a prospective track record. PE funds invest in demonstrated performance, not retrospective reconstruction.

Credibility risk: Medium-High. PE funds will discount retrospective validation heavily.

Remediation: Quantify the expected value of correct WAIT verdicts across the 50-case library. Calculate the aggregate financial impact of the failures the council correctly identified vs. the failures it missed.

Reviewer objection: “Your retrospective accuracy is interesting. But I'm investing in a company that will advise on future decisions. What's your prospective accuracy?”

Documentation Gaps

Gap	Severity	Remediation
No verifiable evidence boundary architecture	Critical	Build document-retrieval corpus with DOI-cited pre-cutoff sources
No published methodology (exact prompts, model version)	Critical	Publish full methodology paper
No inter-run consistency data	High	Run each case twice; report consistency
No independent expert validation of DDIA	High	Commission 3 independent pharmaceutical experts to review council outputs
No pre-specified case selection criteria	High	Publish selection criteria before running expansion cases
No comparator arm	Medium	Design prospective pilot with blinded comparison to internal teams
No confidence intervals on performance metrics	Medium	Calculate and publish CIs for all metrics
No liability framework	Medium	Engage legal counsel on advisory liability

Top 20 Pharmaceutical Executive Objections

1. “Your evidence boundary is a prompt instruction. I can't audit a prompt.”
2. “The LLM was trained on post-2006 data. It knew torcetrapib failed. Your retrospective is contaminated.”
3. “You selected 10 famous cases. Of course you got good results. Show me the cases you didn't include.”
4. “Your 0% false positive rate is based on 5 cases. That's not a rate. That's a coincidence.”

5. "You failed on Verubecestat and then proposed a constitutional amendment to fix it. That's retrofitting, not learning."
6. "The 10 personas are 10 outputs from the same model. That's not a council. That's one opinion repeated 10 times."
7. "I ran your methodology and got a different vote distribution. Which result is correct?"
8. "Your 90% DDIA is self-assessed. Who validated it?"
9. "You have no comparator arm. How do I know a competent internal team wouldn't have caught the same signals?"
10. "You have no prospective track record. Retrospective accuracy is not predictive accuracy."
11. "Your constitutional rules were calibrated on the cases you're using to validate them. That's circular."
12. "You have zero immunology cases in your original library. Your rules were never tested on the most commercially important therapeutic area."
13. "Your financial impact figures are from published sources. The council didn't calculate them. You added them after the fact."
14. "The Verubecestat failure shows your council can't handle Alzheimer's. That's the most expensive therapeutic area in drug development."
15. "Your 'broad alignment' metric is subjective. Who decides what counts as 'identifying the primary decision driver'?"
16. "You have no liability framework. If a company uses your council and the drug fails, who is responsible?"
17. "Your statistical credibility assessment correctly says you can't make statistical claims at 10 cases. So why are you presenting this to us?"
18. "The torcetrapib case is your best case. It's also the most obvious case. Everyone knew the BP signal was a problem. Your council didn't discover anything new."
19. "You have no regulatory acceptance. FDA and EMA have not reviewed or endorsed this methodology."
20. "Your constitution has 10 rules. Real pharmaceutical governance has 200 pages of ICH guidelines. You're not replacing anything."

Remediation Plan

Immediate (before any external presentation):

1. Publish the full methodology: exact prompts, model version, temperature, constitutional rules, persona definitions
2. Disclose the evidence boundary limitation explicitly: "The LLM training data cannot be verified as pre-cutoff. This is a known limitation."
3. Calculate and publish confidence intervals for all performance metrics
4. Commission one independent pharmaceutical expert to review the torcetrapib case output

Before enterprise pilot discussions:

1. Build a document-retrieval architecture with verifiable pre-cutoff evidence corpus
2. Run each of the 10 original cases twice; publish inter-run consistency data
3. Pre-specify the selection criteria for the 40 expansion cases
4. Commission 3 independent experts to validate DDIA for the original 10 cases

Before peer-review submission:

1. Complete 50 cases with pre-specified methodology
2. Design and launch a blinded prospective pilot
3. Engage FDA/EMA for informal methodology review
4. Engage legal counsel on advisory liability framework

DELIVERABLE 8 — Executive Commercialisation Recommendation

The Bet

For AgenThink Mesh to be commercially viable as a pharmaceutical governance tool, one thing must be true: **the council must consistently surface material decision signals that are visible in pre-decision data and that human teams demonstrably missed or discounted.**

The torcetrapib case proves this is possible. The Verubecestat case proves it is not guaranteed. The 40 expansion cases will determine which of these is the rule.

At 10 Cases: NO-GO for Commercialisation

Verdict: NO-GO

The 10-case library is not sufficient for commercialisation. The reasons are not primarily about performance — the 90% DDIA and 0% false positive rate are genuinely encouraging. The reasons are methodological:

The evidence boundary is unverifiable. The methodology is unpublished. The performance metrics have confidence intervals too wide to support any claim. The single failure (Verubecestat) has not been validated by independent experts. The constitution has not been tested on the most commercially important therapeutic areas (immunology, oncology at scale).

What 10 cases supports: Internal demonstration. Proof of concept presentations to sophisticated early adopters. Recruiting pharmaceutical advisors. Designing the prospective pilot.

What 10 cases does not support: Any accuracy claim. Any enterprise contract. Any regulatory discussion. Any investor presentation that implies validated performance.

At 20 Cases: CAUTION — Pilot-Eligible

Verdict: CAUTION

At 20 cases, AgenThink can begin structured pilot discussions with pharmaceutical companies, provided:

1. The methodology is fully published
2. The evidence boundary limitation is explicitly disclosed
3. Performance metrics are presented with confidence intervals
4. The pilot is structured as a research collaboration, not a commercial service
5. No accuracy claims are made beyond “statistically distinguishable from random chance” (if balanced accuracy $\geq 75\%$)

Commercial structure at 20 cases: Research partnership agreements with 2–3 pharmaceutical companies. No fee for service. Data sharing in exchange for access to current Phase II decision data. Prospective pilot design agreed before execution.

At 50 Cases: GO (Conditional)

Verdict: GO (Conditional)

At 50 cases with pre-specified methodology, independent expert validation, and published performance metrics with confidence intervals, AgenThink can:

1. Present to pharmaceutical R&D leadership as a validated decision-support tool
2. Offer enterprise pilots with disclosed limitations
3. Submit a methods paper to a peer-reviewed journal
4. Engage FDA/EMA for informal methodology review
5. Raise Series A financing with a defensible validation story

Conditions for GO:

- Balanced accuracy ≥ 0.75 (95% CI lower bound ≥ 0.60)
- False positive rate $\leq 20\%$ (council does not block more than 1 in 5 successful drugs)
- Independent expert validation of DDIA for at least 20 cases
- Published methodology paper
- Evidence boundary architecture (document retrieval, not prompt instruction)

Commercial structure at 50 cases: Fee-for-service enterprise pilots. 50K–200K per case review. Target: 5–10 pilot clients in Year 1.

At 100 Cases: GO

Verdict: GO

At 100 cases, AgenThink has a peer-review-eligible dataset, statistically credible performance metrics, and a prospective pilot track record (if the pilot was launched at 20 cases, results will be available by 100 cases).

Commercial structure at 100 cases: Enterprise SaaS licensing. Regulatory submission support (with appropriate disclaimers). Strategic partnership with a Big Four advisory firm or CRO. Series B financing.

The Honest Summary

Milestone	Verdict	Rationale
10 cases	NO-GO	Methodology unverifiable; metrics too wide; no independent validation
20 cases	CAUTION	Pilot-eligible with full disclosure; no commercial claims
50 cases	GO (conditional)	Enterprise pilots defensible; peer review eligible
100 cases	GO	Full commercialisation; prospective track record

The most important single action AgenThink can take today is not running more cases. It is building the verifiable evidence boundary architecture. Without it, every case in the library is vulnerable to the objection: “The model knew the answer.” That objection, if unanswered, will end every serious conversation with every serious reviewer.

Run the 40 expansion cases. But build the architecture first.

APPENDIX — Proposed 50-Case Execution Sequence

The following sequence is designed to build the validation track record in the most defensible order: starting with the highest-governance-value cases and building toward the clean-success stress tests.

Batch	Cases	Rationale
Batch 2 (Cases 11–20)	Oncology 10	Tests false-positive blocking on accelerated approval; high public evidence
Batch 3 (Cases 21–28)	Cardiovascular 8	Tests cangrelor persistence, vorapaxar benefit/risk, mavacamten novel mechanism
Batch 4 (Cases 29–34)	Neurology/AD 6	Validates PC-011 amendment; tests gantenerumab and crenezumab
Batch 5 (Cases 35–42)	Immunology 8	Primary false-positive stress test; all 8 drugs succeeded
Batch 6 (Cases 43–48)	Metabolic 6	Tests SGLT2 class (empagliflozin CV surprise), GLP-1 obesity
Batch 7 (Cases 49–53)	Rare Disease 5	Tests gene therapy evidence thresholds; first-in-class mechanisms

After each batch: Update the metrics dashboard. Assess whether a constitutional amendment is warranted. Publish the batch results before proceeding to the next batch.